

# Assignment 1: Gen AI Application

## Problem Statement

You are required to build a **Retrieval-Augmented Generation (RAG)** application with **agentic behavior**, a **Streamlit UI**, and a **FastAPI backend**.

### Your tasks:

#### 1. Create a Knowledge Base

- Use any publicly available dataset (e.g., airline SOP documents, financial policies, retail product descriptions, etc.).
  - Convert documents into embeddings and store them locally (e.g., FAISS, ChromaDB).
- 

#### 2. Build a RAG Pipeline (in FastAPI)

Your FastAPI application must expose the following endpoints:

##### 1. /ask

- Input: { "query": "..." }
- Output: A GenAI response grounded in retrieved documents.

##### 2. /agent-task

- Implement a simple **task-solving agent** that:
    - Breaks a user query into smaller sub-queries
    - Retrieves relevant context
    - Produces a final synthesized answer
  - Example queries the agent should handle:
    - “Compare policies between Document A and B and summarize differences.”
    - “Find steps from all SOPs and create a combined checklist.”
- 

#### 3. Streamlit Application

Build a frontend with:

- A text input box with chat interface

- A dropdown to select **RAG** or **Agent Mode**
  - Display retrieved context + final answer
  - Call your FastAPI backend
- 

#### 4. Deployment

You must deploy the app so we can test it:

- Provide:
    - Public URL for testing
- 

#### 5. Submission Requirements

- Submit a **ZIP file of your codebase**, containing:
  - FastAPI backend code
  - Streamlit UI
  - RAG pipeline
  - Agent logic
  - Requirements.txt
- A video explaining the working of the GenAI application

## Assignment 2: ML Lifecycle

### Problem Statement

Consider any open dataset -

**Example:** Telco Customer Churn dataset, Loan Default dataset, or any Classification dataset.

Your task:

#### 1. Perform EDA

- Identify missing values, outliers
- Understand feature distributions
- Identify key drivers of the target variable

- Provide visualizations
  - Summarize insights in business terms
- 

## 2. Preprocessing & Feature Engineering

- Handle missing values
  - Encode categorical features
  - Scale numerical variables if needed
  - Engineer at least **2 new features**
- 

## 3. Build and Train Models

Train at least **three different models**, e.g.:

- Logistic Regression
  - Random Forest
  - XGBoost / LightGBM
- 

## 4. Evaluate Model Performance

Include:

- Accuracy
  - Precision, Recall, F1
  - ROC curve & AUC
  - Confusion matrix
  - Feature importance
  - Business interpretation of results
- 

## 5. Select the Best Model

Explain:

- Why this model performs best
- How it would be deployed

- What monitoring would look like in production
- 

## **6. Deliverables**

Submit a ZIP containing:

- Jupyter notebook
- Dataset (or link)
- Model files
- A video explaining the solution & the full life cycle